- Evaluating NLP systems on individual tasks seems a bit silly. A better approach to evaluating a model's language understanding capability is to evaluate it on many tasks. This is the idea behind the GLUE benchmark.

The GLUE benchmark

24

GLUE Leaderboard

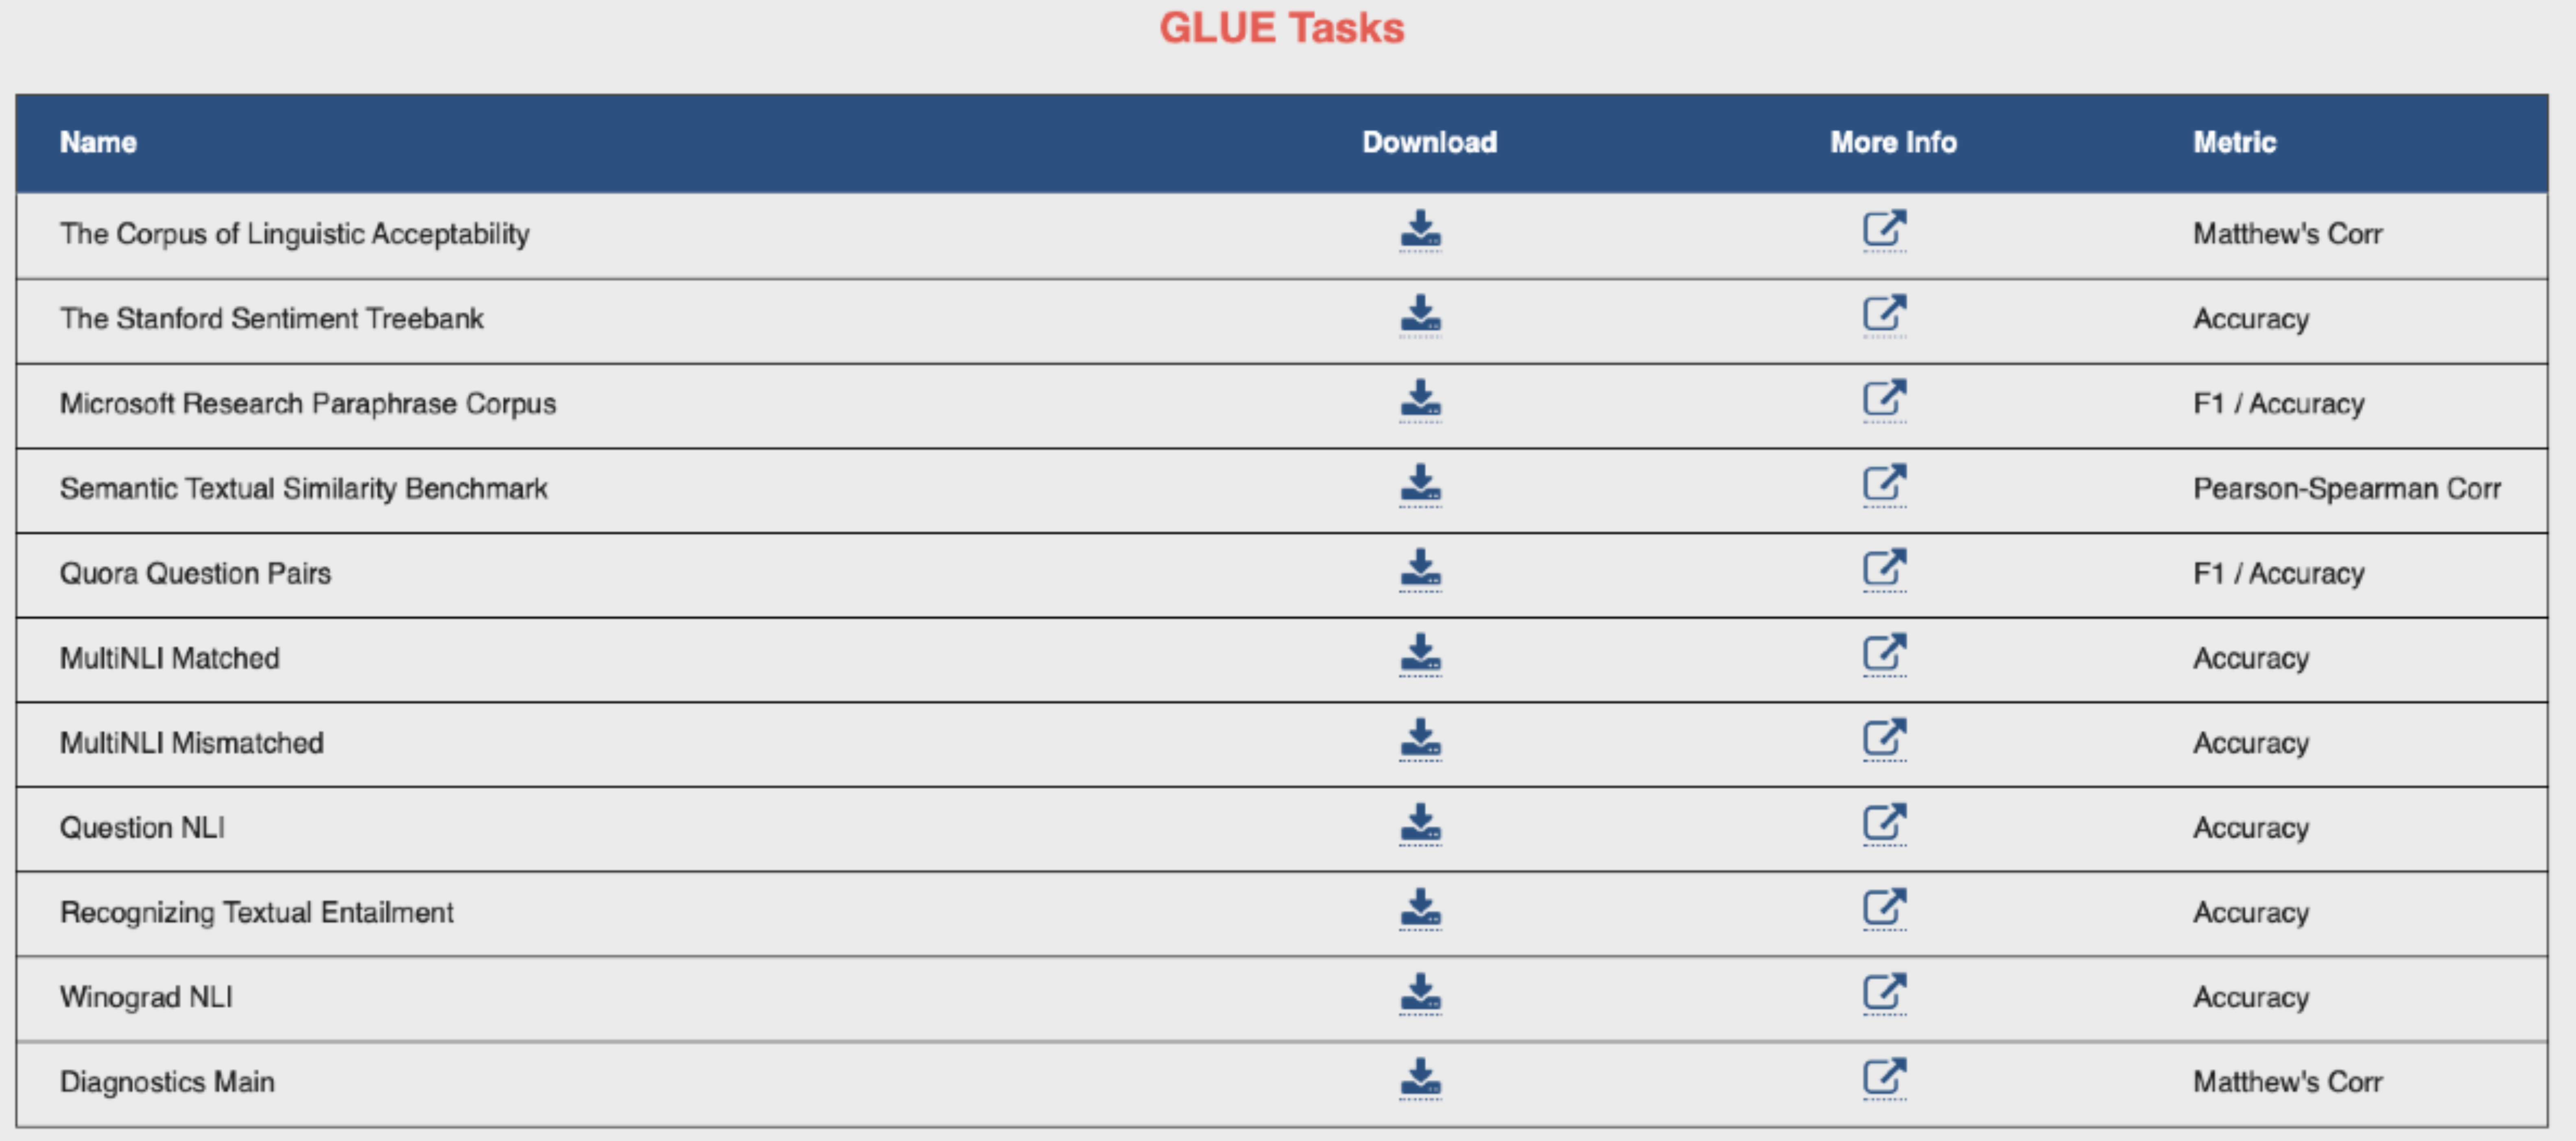

GLUE Tasks

Name	Download	More Info	Metric
The Corpus of Linguistic Acceptability	⬇	↗	Matthew's Corr
The Stanford Sentiment Treebank	⬇	↗	Accuracy
Microsoft Research Paraphrase Corpus	⬇	↗	F1 / Accuracy
Semantic Textual Similarity Benchmark	⬇	↗	Pearson-Spearman Corr
Quora Question Pairs	⬇	↗	F1 / Accuracy
MultiNLI Matched	⬇	↗	Accuracy
MultiNLI Mismatched	⬇	↗	Accuracy
Question NLI	⬇	↗	Accuracy
Recognizing Textual Entailment	⬇	↗	Accuracy
Winograd NLI	⬇	↗	Accuracy
Diagnostics Main	⬇	↗	Matthew's Corr

The GLUE benchmark

- Evaluating NLP systems on individual tasks seems a bit silly. A better approach to evaluating a model's language understanding capability is to evaluate it on many tasks. This is the idea behind the GLUE benchmark.

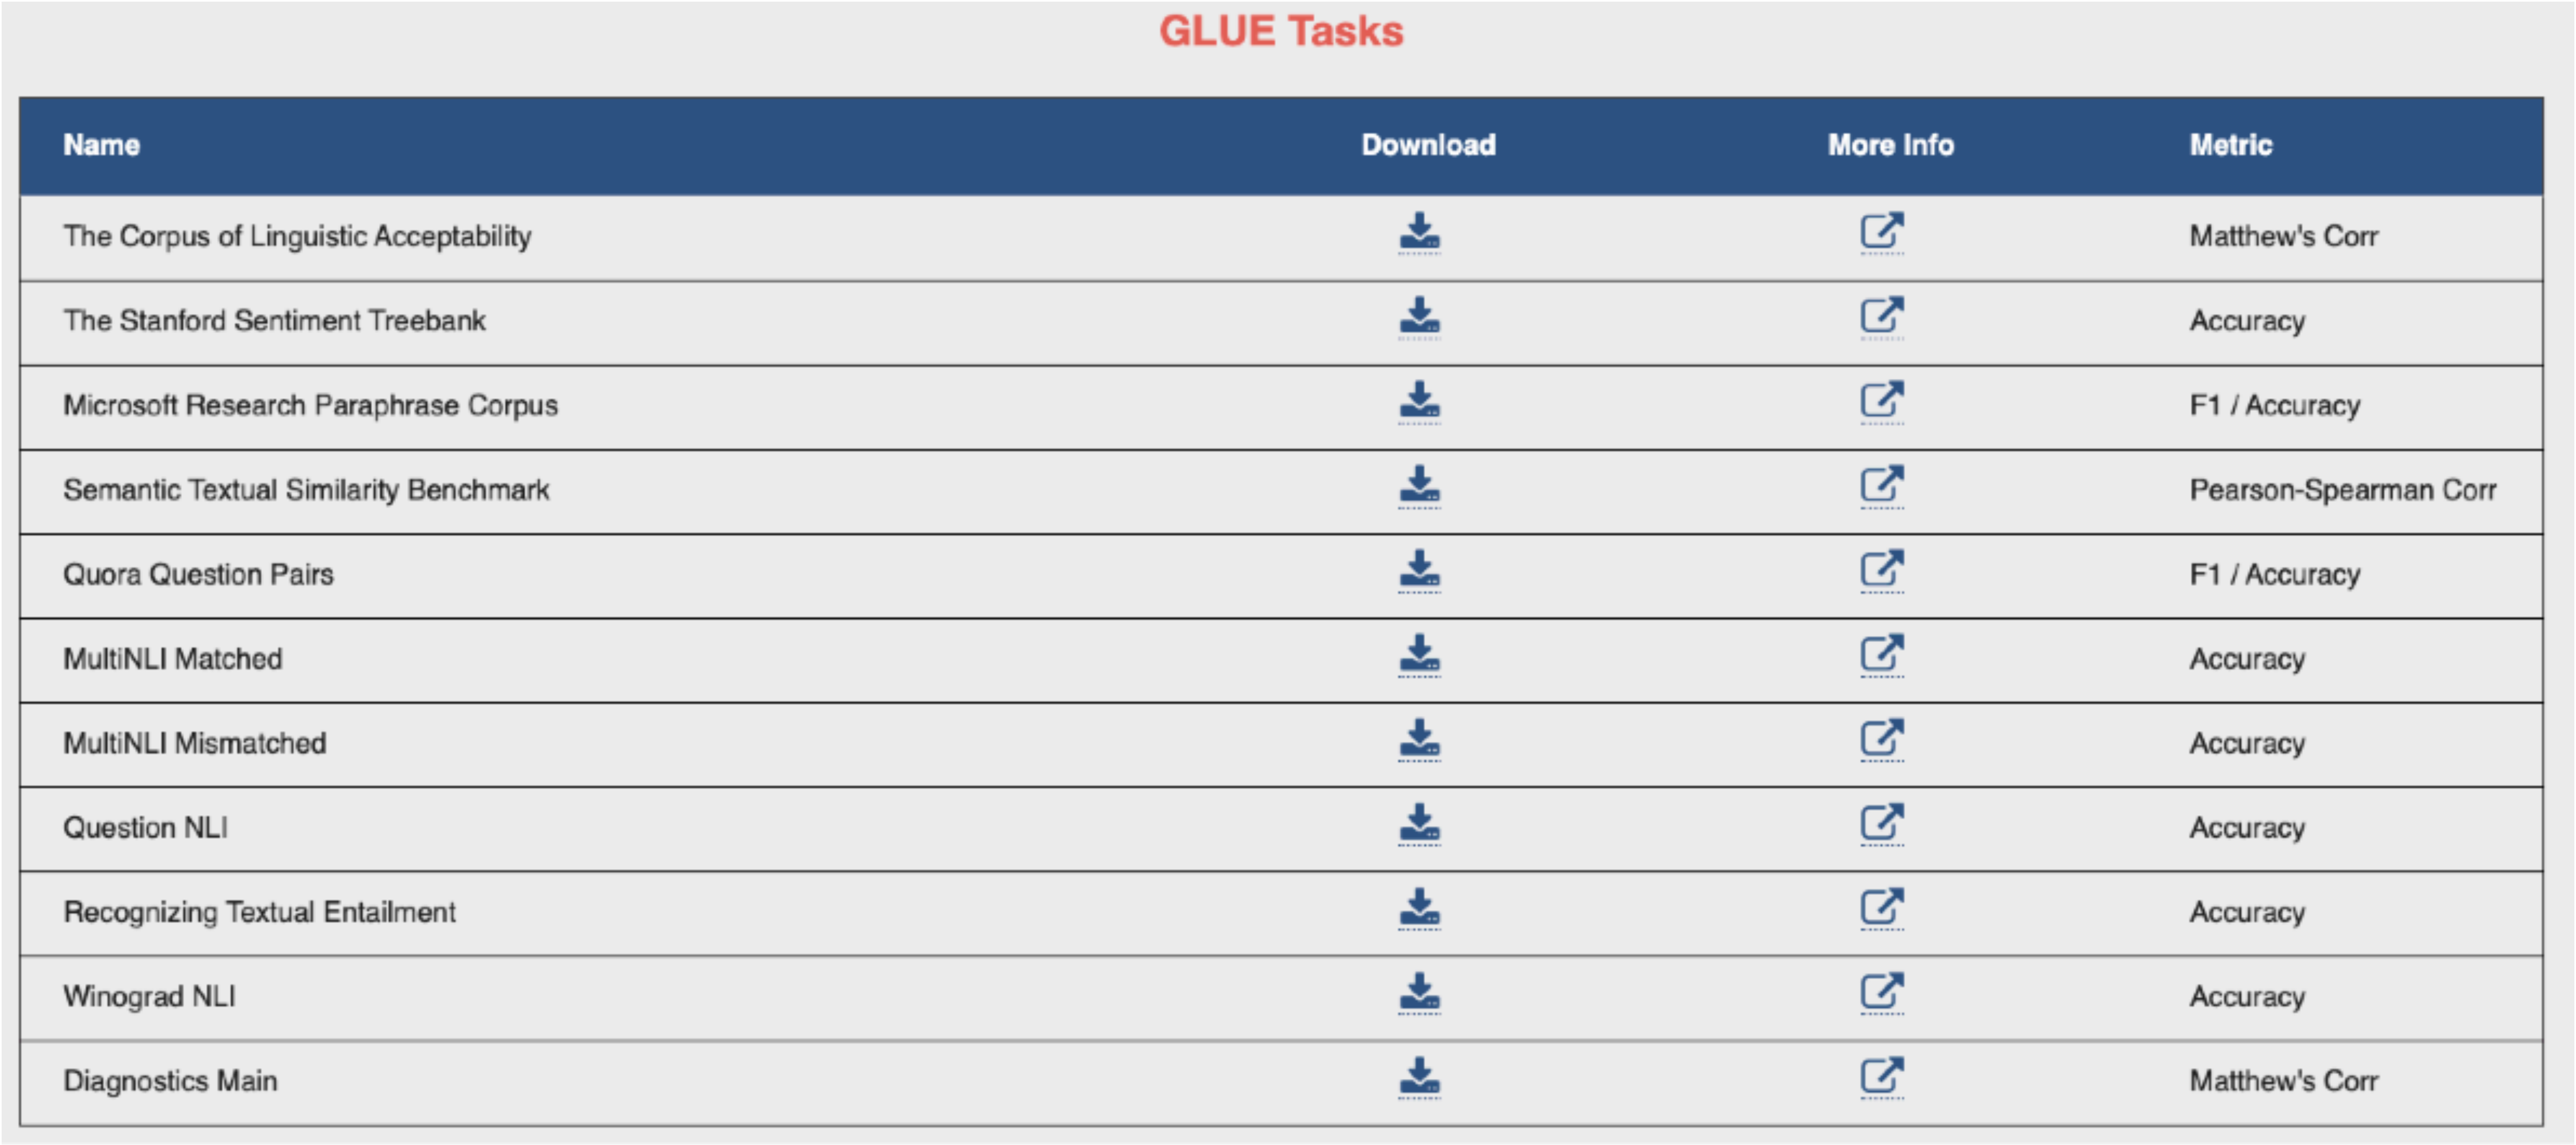

GLUE Tasks

Name	Download	More Info	Metric
The Corpus of Linguistic Acceptability	⬇	↗	Matthew's Corr
The Stanford Sentiment Treebank	⬇	↗	Accuracy
Microsoft Research Paraphrase Corpus	⬇	↗	F1 / Accuracy
Semantic Textual Similarity Benchmark	⬇	↗	Pearson-Spearman Corr
Quora Question Pairs	⬇	↗	F1 / Accuracy
MultiNLI Matched	⬇	↗	Accuracy
MultiNLI Mismatched	⬇	↗	Accuracy
Question NLI	⬇	↗	Accuracy
Recognizing Textual Entailment	⬇	↗	Accuracy
Winograd NLI	⬇	↗	Accuracy
Diagnostics Main	⬇	↗	Matthew's Corr

GLUE Leaderboard

Lab 07: Project Kickoff